

LINK QUALITY ESTIMATOR FOR A MOBILE ROBOT

Narek Pezeshkian, Joseph D. Neff, and Abraham Hart
SPAWAR Systems Center, Pacific, San Diego, CA 92152
narek.pezeshkian@navy.mil

Keywords: Link Quality: Video Quality: Estimator: Metric: Mobile Robot

Abstract: Maintaining link connectivity between a mobile robot and its control station in a non-line-of-sight environment is challenging. One solution is to use intermediate relay radios that the robot can carry and deploy when and where needed to maintain the link. However, the precise placement locations for the relays are not known ahead of time. Therefore, the deployment decision must be formulated online and the relays deployed before the link with the control station breaks. A link-quality estimator is developed based on video throughput and received signal strength indicator data. The estimator takes into account human perception of video quality that is obtained via subjective testing by an operator. The data is used to train the link-quality estimator, which issues an alert that can be used as a trigger for an automatic relay deployment mechanism or to advise the operator to manually deploy relays before the link between the robot and control station fails.

1 INTRODUCTION

Tactical mobile robots have been increasingly used by the military over the past several years. This is especially true for Explosive Ordnance Disposal (EOD) teams that use robots to investigate and neutralize Improvised Explosive Devices. These robots are remotely controlled from the operator control unit (OCU) using digital radios. The high operating frequency of these radios requires a line-of-sight (LOS) to the OCU, which is difficult to maintain in urban environments. The link between the robot and OCU can fail, usually rather quickly, when operating beyond LOS due to multipath interference and signal fading.

Controlling a robot via a tethered connection, typically fiber-optic, eliminates the LOS problem but introduces new ones. Tethered connections can snag and break, limiting mobility. Advanced radio systems that utilize sophisticated modulation techniques and take advantage of MIMO antenna technology thrive in multi-path environments and can overcome the LOS limitations to a degree. However, obstacles that severely block and attenuate the signal can still be problematic.

The use of relays, on the other hand, adds an unprecedented degree of freedom to where robots can operate. Relays can entirely overcome severe

obstacle blockages so long as a LOS can be maintained with adjacent radios in a chain of relays. Determining the placement location of such relays is critical. The focus of this paper is the formulation of a link-quality (LQ) estimator, the output of which is used either by the robot (automatically) or the operator (command sent from the OCU) to release a relay before the link breaks. Section 2 provides a brief background of various relay systems designed for tactical robots. Section 3 discusses the LQ estimator design. Simulation results are outlined in section 4, and section 5 concludes the paper.

2 BACKGROUND

The solution to address the LOS requirement between a tactical robot and its OCU began in 2002 under the Autonomous Mobile Communications Relay (AMCR) project (Nguyen, H. G., Pezeshkian, N., Raymond, M., Gupta, A., Spector J. M., 2003). The goal of the AMCR system was to provide extended range and non-line-of-sight (NLOS) operational capability for tactical robots. This was accomplished through the use of dedicated mobile relay robots (or mobile nodes) that followed the lead robot in a convoy formation and automatically stopped when needed to maintain the link. The

radios on-board the lead robot, mobile nodes, and OCU formed a mesh network that allowed the operator to teleoperate the lead robot based on video relayed to the OCU.

The mobile nodes must be set up in a specific order in such a convoy since each mobile node is programmed to follow the robot in front of it. In addition, each mobile node is programmed to monitor the received signal strength indicator (RSSI) data of the node immediately behind it. For example, the last mobile node in the convoy monitors the RSSI of the OCU. The RSSI data, which is used as a measure of the link quality, is compared to a predetermined threshold, below which the mobile node stops to maintain the link.

The AMCR solution proved to be very successful and the commercial-off-the-shelf (COTS) 802.11b radios and processor boards were extremely cost effective. However, the AMCR system was a research project and never designed for field use, since the mobile nodes were expensive and logistically impractical.

A more realistic solution was developed under the Automatically Deployed Communication Relays (ADCR) project (Pezeshkian, N., Nguyen, H. G., Burmeister, A., 2007). The ADCR system shown in Figure 1 consists of a Deployer and several Relay “Bricks”. The Deployer carries the Relay Bricks and mounts onto a small ground robot.



Figure 1: ADCR Deployer mounted on an iRobot *PackBot* with one deployed and five stowed Relay Bricks.

The Deployer and the Relay Bricks each have the same radio hardware and RSSI-based link-quality estimator that is used by the AMCR system. However, the only mobile node is the Deployer, therefore, the link monitoring and the decision to eject a Relay Brick is formulated by the Deployer radio. Once a Relay Brick is ejected it self-rights and extends the antenna. As the operator controls

the robot along its path more Relay Bricks are ejected as needed to maintain the link.

The success of ADCR led to several licensing agreements with commercial developers. Subsequent projects led to additional developments that improved upon the system. For example, the redesigned Deployer of the Automatic Payload Deployment System (APDS) (Pezeshkian, N., Nguyen, H. G., Burmeister, A., Holz, K., Hart, A., 2010) allows a robot to carry and deploy not only Relay Bricks but a wide range of other types of payloads, such as leave-behind sensors and containers. The container payloads can be used to carry food, ammunition, medical kits, and anything else that fits within. The Relay Bricks were also redesigned to contain faster radios and an improved antenna lift mechanism as shown in Figure 2.




Figure 2: APDS Deployer mounted on an iRobot *PackBot*. Various payload types are shown around the robot.

The interest that was generated by numerous publications and successful demonstrations of the APDS and ADCR systems led the Naval EOD Technology Division (NAVEODTECHDIV) to fund the development of a robust radio repeater solution for use by currently-fielded robotic vehicles. It was necessary to deliver this solution quickly; therefore, a simplified, albeit robust system based on APDS technology was developed under the Manually Deployed Communication Relays (MDCR) project. The MDCR system omits the RSSI monitoring and automatic-deployment capability of APDS, and instead relies entirely on remote commands sent from the OCU to deploy the Relays as the operator sees fit. Although simple in design, the MDCR system has been successfully field tested with plans to mass produce additional units.

Although RSSI-based link monitoring has been successful in the ADCR and APDS systems, it is not an ideal solution, as will be explained in section 3. Therefore, the goal of the LQ estimator outlined in this paper is to provide a superior estimation method

that will assist the MDCR operator in placing relays, and to also provide a trigger to automatically eject relays for future ADCR systems.

3 QUANTIFYING LINK

In the MDCR system the Relays are deployed based on operator command. Two factors play a role in the Relay deployment decision-making process of the operator: 1) prior knowledge of LOS loss – the operator knows that controlling the robot around a large obstruction will cause a loss of LOS so a Relay is deployed before proceeding, and 2) video degradation – as the distance between the robot and OCU increases, even under LOS conditions, the operator deploys a Relay when video quality degrades.

Although these factors can be effective for deploying Relays, in order to maintain the link between the robot and the OCU, the operator for the most part is guessing as to where to place the Relays based on experience and intuition about the RF environment. If the relaying system could provide an indicator based on some sort of LQ estimator that can warn of a failing link, however, the operator would be in a much better position to optimize Relay placement. This is important since the number of Relays carried by a robot is limited and maximizing the distance between the Relays translates into maximizing the stand-off distance of the robot. Furthermore, the LQ estimator can be used by a relaying system (e.g., ADCR) to provide automatic Relay deployment capability, effectively alleviating the operator from the deployment task.

It is also important to keep in mind that the link under consideration is between the robot and the next-hop neighbor of the routing path leading back to the OCU. This is, in fact, the only dynamic link given that the only mobile node is the robot and all other nodes (OCU and previously deployed Relays) are static.

3.1 Link Quality

In this section a background on recent work on link quality is given, followed by sections that describe the proposed LQ metrics used by the LQ estimator.

3.1.1 LQ Background

A plethora of research on LQ estimation can be found in the literature. Many schemes combine multiple variables available from the physical and

link layers to form a more comprehensive and robust LQ metric. Rondinone, Ansari, Riihijärvi, and Mähönen (2008) propose multiplying the Packet Reception Rate (PRR) of a link by the corresponding mean RSSI value to obtain a new LQ indicator that can be used by a network to select an optimal routing path. Srinivasan, Kazandjieva, Jain, and Levis (2008) combine PRR and channel burstiness to estimate TCP throughput. Liu and Cerpa (2011) combine RSSI, PRR, signal-to-noise ratio (SNR) and the Link Quality Indicator (LQI) provided by the CC2420 radio chip to provide a probability of successfully delivering the next packet.

Yet combining variables is not the only approach. Farkas, Hossmann, Ruf, and Plattner (2006) propose using pattern matching to predict the future behaviour of a link. Each node keeps a time series record of the SNR with each of its links and uses pattern matching to find the best match in an attempt to estimate the future behaviour of the SNR. Qin, He, and Voigt (2011) develop a new LQ estimator, called the Spectrum Factor (SF), which is derived from frequency-domain data.

3.1.2 LQ Data

An LQ estimator can be used by a routing protocol in a mesh network to select optimal routing paths (Liu, Fan, Shu, Yu, 2010 and Liu and Cerpa, 2011). The goal of the LQ estimator for the MDCR system is somewhat different: Develop an LQ estimator that is suitable in predicting link failure such that a Relay can be deployed before the link breaks.

The LQ estimators discussed in the previous section are unsuitable for use given the stated goal. Rondinone et al. (2008) suggest multiplying the PRR of a link by the corresponding mean RSSI value to help in selecting routing paths. Since there is only one link under consideration (between robot and next-hop neighbor along the routing path leading to the OCU), this multiplication provides no new information. Srinivasan et al. (2008) attempt to estimate TCP throughput, which is unnecessary since the video data of the robot uses UDP packets and the throughput is readily available. Liu et al. (2011) make use of SNR and LQI data that is unavailable in the 802.11 radios used in the MDCR system. Farkas et al. (2006) use pattern matching to predict future behaviour of a link. This requires some level of repetitive pattern to be present in the collected data, which is highly unlikely given the random movements of a teleoperated robot. Finally, Qin et al. (2011) estimate LQ in the frequency

domain, which requires raw RF data that is not easily obtainable from the MDCR radios.

The data selected for the development of the proposed LQ estimator is UDP throughput (packets-per-second) and RSSI, which are readily available and ease integration of the estimator into the existing mesh network software of the MDCR system. The throughput data is also a direct indicator of video quality – one of the key factors in the deployment decision-making process of the operator. Video quality, however, is subjective. A slightly choppy video may be acceptable to one operator and unacceptable to another. To quantify video quality, an experiment was devised where an operator controlled the robot along a predetermined path and when the video quality, as judged by the operator, began to degrade, the operator marked that point in time. The marking method is simply a key press on a test laptop that collects throughput and RSSI data along with operator key presses, all synchronized in time. There were two different key presses involved in this experiment: The #2 key was pressed when video quality began to degrade and the #3 key was pressed when the link was completely lost. These two moments in time are t_F (failing) and t_L (lost), respectively. The link-failure period (t_{LF}) is simply $t_L - t_F$. A sample of collected data and key presses is shown in Figure 3. A simple moving average (MA) process is applied to all data to smooth out variations.




Figure 3: Example of video throughput (blue) and RSSI data (green) received at the OCU from a *PackBot* using the MDCR system. Solid line is the average (μ) of past five samples of underlying (dotted) data. Left and right black lines represent t_F and t_L , respectively.

Many such trial runs were performed under two different environments, one more prone to multipath than the other. In all test trials, clear trends are observed in the throughput data during t_{LF} ,

summarized as follows: 1) The throughput begins to roll off sometimes gradually and sometimes relatively sharp, and 2) the throughput variance increases. The RSSI data, as expected, drops gradually overtime as the robot moves away from the OCU. Before t_F , however, the throughput data does not show any clear trend. The test trials show that t_{LF} varies between 10 to 20 seconds, which provides ample time to issue an alert. These trends are exploited in the design of the LQ estimator.

3.1.2 RSSI as Early Warning

RSSI data has been proposed as a good link-quality metric by Srinivasan and Levis (2006) but the limitations of this statement must be understood. It has been shown by Vlavianos, Law, Broustis, Krishnamurthy, and Faloutsos (2008) that RSSI data is measured at the lowest rate and cannot characterize the LQ at high transmission rates. Furthermore, RSSI is only measured from the packet preamble; therefore, if an interfering signal happens to prevent proper reception of the preamble, the RSSI will simply not be recorded. If the interfering signal happens to corrupt the packet after the preamble has been received, then the RSSI will be recorded as if there is no interferer. Hence, RSSI data is unchanged even in the presence of an interferer. The work of Judd, Wang, and Steenkiste (2008) further supports this assessment.

Broadband noise, however, is a concern. If the overall noise floor is raised due to external broadband sources of noise, the overall SNR of received packets will decrease. This means that RSSI data can only be measured down to the raised noise floor since packets received below the noise level will be corrupted. Looking at Figure 3 it may seem reasonable to threshold the RSSI at about 10, below which the throughput data enters the region of degraded video quality, t_{LF} . This approach may work in the absence of broadband noise, but that constraint cannot be guaranteed when operating in a variety of environments.

The goal of the proposed LQ estimator is to predict link failures so that a Relay can be deployed before the link breaks. Preferably, some early warning should be given to the operator by the relaying system, followed by an imminent failure alert so that the operator can deploy a Relay before the link breaks. Interfering signals are not a major concern given the operating environment, where the overwhelming reason for link failure is due to signal fading and loss of LOS. Broadband noise, however, can exist. Given the limitations of RSSI, it is then

reasonable to use it only as a conservative early warning system. Figure 4 shows the mean RSSI value exactly at time t_F , for all test trials. There are clear variations but the overall range is low. A conservative early warning of link failure can be issued, for example, if the mean RSSI drops below a threshold of 20. Selecting a high threshold leaves quite a bit of margin should the noise floor increase due to broadband noise.




Figure 4: Mean RSSI value at time t_F for all test trials.

3.1.3 Link Quality Metrics

An accurate estimation of imminent link failure is required to alert the operator of complete loss of connectivity or trigger the deployment of a Relay from an automated deployment system. Since RSSI data does not accurately reflect the ability of a link to successfully deliver packets as discussed in section 3.1.2, throughput data is used instead.

The first trend of the throughput data is increased variance σ during t_{LF} . At the same time, the mean μ drops due to the second trend, the roll off. Since the mean is high and variance low prior to t_F and vice versa during t_{LF} , it is reasonable then to use the ratio of the two as a metric. This is inspired from the Ricean K-factor (Greenstein, Michelson, and Erceg, 1999), which is used as a measure of signal fading. The ratio here is given as $\kappa = \mu/\sigma$ and is the first LQ metric, LQM_κ .

The second trend is the roll off. This is measured by first taking N samples of throughput data then calculating its intercept (x_1) and slope (x_2) using linear regression. The assumption is that the N -sample-long data is a straight line. Using a sliding window, x_1 and x_2 are updated for each new sample. The vector $\mathbf{x} = [x_1 \ x_2]$ is the second LQ metric, LQM_x .

The trade-off between the false-alarm rate and the miss rate is dependent on N . Low false-alarm

and low miss rates are desired. By setting N too high, the data will be too smooth and the LQ estimator slow to respond. This has the effect of reducing the false-alarm rate due to reduced noise, but increases the miss rate due to reduced response time. In effect, the link is lost before the LQ estimator has a chance to issue an alert. On the other hand, setting N too low causes the data to be too noisy, increasing the false-alarm rate, but reducing the miss rate due to increased response time. Since the cost of failing to issue an alert (a miss) is much greater than alerting too soon (a false alarm), the selection is biased towards reducing the miss-rate by choosing $N = 5$.

3.1.4 Classifier

The keystrokes of the operator during the test trials essentially label the collected data that are used to train the LQ estimator. Half of the collected data is used as training data and the other half as test data. A labelling problem can be solved by classifiers. Supervised training is used by two classifiers, one for LQM_κ and the other for LQM_x . Each classifier finds the optimal decision boundary between two different sets of labelled data: those marked before t_F (signal OK) and those marked during t_{LF} (signal failing). The hypothesis function for LQM_κ is given by $z_\kappa(\theta_\kappa) = \theta_0 + \theta_1\kappa$ and for LQM_x the hypothesis function is $z_x(\theta_x) = \theta_0 + \theta_1x_1 + \theta_2x_2$. Both are modelled as linear functions, which is a reasonable assumption when looking at the data clusters in Figures 5 and 6. The optimal parameter vector θ is found by the classifier, which defines the decision boundary that has values $z(\theta) \geq 0$ on one side and $z(\theta) < 0$ on the other.

The plot of the labelled κ values for all test trials is shown in Figure 5. The plot of labelled x_1 and x_2 values for all test trials is shown in Figure 6. The green circles represent values that take place before t_F and the red asterisks are data that take place during t_{LF} . Using logistic regression, an optimal decision boundary is generated, shown as the blue line. All green circles above the line are hits (link OK) and those below the line are false alarms (link failing when in fact it is not). All red asterisks below the line are hits (link failing) and those above the line are misses (link failing but no alert issued). It is clear from both figures that there is overlap between the labelled data. Given the high cost of misses, the decision boundary is biased so as to reduce the number of misses.

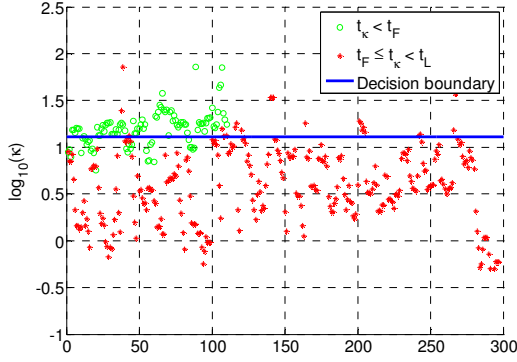


Figure 5: Plot of labelled κ and decision boundary. Its log is taken to improve computation of the boundary.

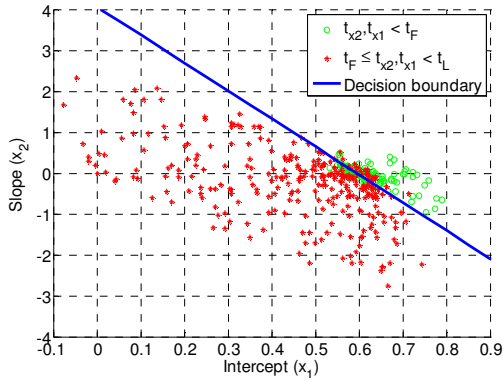


Figure 6: Plot of x_1 vs. x_2 . Both variables have been scaled to reduce their range for improved computation of the decision boundary.

Figure 6 supports the roll-off trend of the throughput data. Looking at Figure 3, the flat part of the throughput data roughly corresponds to 40 packets-per-second and since it is flat its slope is about zero. This correlates to the green cluster seen in Figure 6. As the throughput begins to fail during t_{LF} , the packet rate drops, which corresponding to the reduced x_1 (intercept) values. At the same time the slope increases in the negative direction. This corresponds to the red cluster in Figure 6. The positive x_2 values are due to the variance of the throughput data during t_{LF} that can cause the slope to go positive momentarily. Notice, however, very few occurrences of positive slope and high intercept values take place. The occurrences of high intercept and high negative slope can be explained by sharp roll-offs, where the throughput value is still somewhat high but the slope is steep.

3.1.5 Link-Quality Estimator

The goal of the LQ estimator is to provide an early warning of link failure (based on RSSI data) and a more accurate imminent link-failure alert (based on LQ metrics calculated from throughput data). These metrics are somewhat noisy due to the selection of N chosen to increase responsiveness (reduced miss rate), and hence, sensitivity (increased false-alarm rate). Each metric alone is not sufficient to provide an accurate estimation, therefore they are combined. The manner in which they are combined is essentially an AND operation between the hypothesis functions. This implies that both hypothesis functions $z_{\kappa}(\theta_{\kappa})$ and $z_x(\theta_x)$ must agree that the link is failing, which occurs when both $z_{\kappa}(\theta_{\kappa})$ and $z_x(\theta_x)$ are less than zero. Furthermore, the LQ estimator does not issue an alert unless both $z_{\kappa}(\theta_{\kappa})$ and $z_x(\theta_x)$ are less than zero for three consecutive samples in a row. This eliminates momentary glitches where both hypothesis functions are below zero. Finally, the LQ estimator does not start calculating the imminent link-failure alert until a warning is issued when the mean RSSI data falls below a conservative threshold. A simplified flow chart for the LQ estimator is shown in Figure 7.

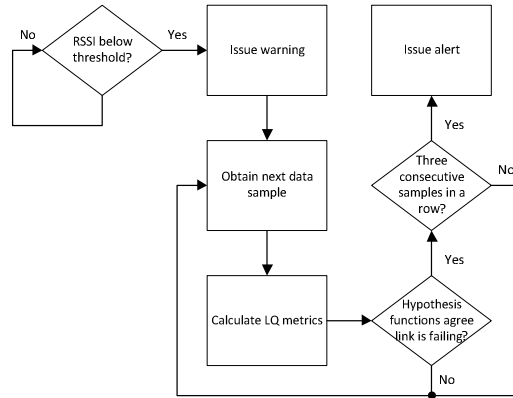


Figure 7: Simplified flow chart for the LQ estimator.

The flow chart does not show the additional steps taken to deactivate the warning and alert indicators. For example, instead of using a single threshold, hysteresis can be added to the mean RSSI data where falling below the lower threshold (e.g., robot moving away from OCU) causes a warning to be issued, which is removed when the mean RSSI moves above the upper threshold (e.g., robot moving back towards OCU). In a somewhat similar manner the link-failure alert indicator can be removed. For example, an issued alert can be removed if both

hypothesis functions agree that the signal is good, say for five consecutive samples.

Figure 8 shows a sample of a test trial. The plot shows that a warning is issued when the mean RSSI falls below 21 (hysteresis enabled). Once the warning has been issued, the LQ estimator begins calculating the LQ metrics and testing the hypothesis functions $z_k(\theta_k)$ and $z_k(\theta_k)$. A value of less than zero indicates a hit, which is shown on the plot as a red box for LQM_k and a red diamond for LQM_k . A link-failure is indicated as a red '+' sign when both functions are less than zero at the same time (LQM_{AND}). An occurrence of this takes place at time 190 but no alert is issued. The LQ estimator issues an alert when it observes three consecutive link-failure hits at time $t_A = 202$. This occurs just after $t_F = 200$, with plenty of time still left before the link is completely lost at time $t_L = 213$.

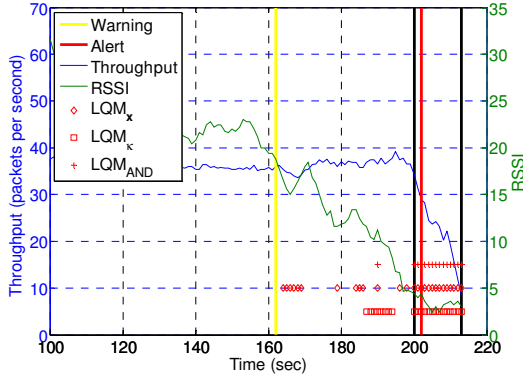


Figure 8: LQ estimator warning and alert. A warning is issued based on RSSI and an alert based on throughput.

4 SIMULATION RESULTS

The LQ estimator will occasionally issue an alert prior to t_F due to the overlap in the training data as shown in Figures 5 and 6. This is a desirable effect because the alert is issued just before video degradation begins. An alert issued after t_F is also acceptable so long as the alert does not take place too close to t_L , which may not provide enough time to deploy a Relay before the link breaks. Therefore, the accuracy of the LQ estimator is defined as the percentage of alerts issued within a specified window of time t_w centered on t_F for all test data. The window t_w is defined as $t_F \pm \Delta t$. The value Δt is equal to βt_{LF} where $0 < \beta \leq 1$. This ensures that Δt is no greater than t_{LF} . The selection of β is somewhat arbitrary. The smaller it is, the closer the alert issue-

time t_A must be to t_F before the alert is counted as an accurate hit. Table 1 shows the accuracy result for different values of β .

Table 1: LQ estimator accuracy

β	Hit %	False Alarm %	Miss %
1/2	73	9	18
2/3	82	9	9
3/4	91	0	9
1	100	0	0

Table 1 shows that with $\beta = 1$ all alerts are issued within $t_F \pm t_{LF}$, and 73% of alerts are issued within $t_F \pm t_{LF}/2$ with $\beta = 1/2$.

5 CONCLUSION

A link-quality (LQ) estimator is developed to provide an accurate means of estimating an imminent link failure, which is required to assist the operator of a tactical mobile robot in deploying a Relay before the link breaks. Since the robot carries a limited number of Relays, increasing the distance between deployment locations will increase the operational range of the robot. The same LQ estimator can also be used on an automatic Relay deployment mechanism (such as the ADCR system) as a trigger to eject a Relay.

The LQ estimator is based on LQ metrics calculated from labelled throughput data. The throughput data is labelled during test trials by the robot operator, who marks the data when the video quality begins to degrade and finally lost altogether. This process is repeated for several trial runs in two different operating environments. The labelled data is used to train the LQ estimator, which is then applied to test data that is not used in the training session. The LQ estimator issues two alerts: 1) a warning alert to the operator based on RSSI data, which serves as a conservative estimate of a link beginning to fail, and 2) a much more accurate link-failure alert based on throughput data when an imminent link failure is detected. The results from the test data show that the LQ estimator achieves high accuracy in issuing an alert before the link is completely lost.

REFERENCES

- Farkas, K., Hossmann, T., Ruf, L., Plattner, B., 2006. Pattern Matching Based Link Quality Prediction in Wireless Mobile Ad Hoc Networks. *MSWiM'06*.
- Greenstein, L. J., Michelson, D. G., Erceg, V., 1999. Moment-Method Estimation of the Ricean K-Factor. *IEEE Communications Letter*, 175-176.
- Judd, G., Wang, X., Steenkiste, P. 2008. Efficient Channel-Aware Rate Adaptation in Dynamic Environments. *MobiSys'08*.
- Liu, T., Cerpa, A.E., 2011. Foresee (4C): Wireless Link Prediction using Link Features. *10th International Conference on Information Processing in Sensor Networks*, 294-305.
- Liu, L., Fan, Y., Shu, J., Yu, K., 2010. A Link Quality Prediction Mechanism for WSNs Based on Time Series Model. *Ubiquitous Intelligence & Computing and 7th International Conference on Autonomic & Trusted Computing*, 175-179.
- Nguyen, H. G., Pezeshkian, N., Raymond, M., Gupta, A., Spector J. M., 2003. Autonomous Communication Relays for Tactical Robots. *11th International Conference on Advanced Robotics*, 35-40.
- Pezeshkian, N., Nguyen, H. G., Burmeister, A., 2007. Unmanned Ground Vehicle Radio Relay Deployment System for Non-Line-of-Sight Operations. *13th International Conference on Robotics & Applications*.
- Pezeshkian, N., Nguyen, H. G., Burmeister, A., Holz, K., Hart, A., 2010. *A Modular Design Approach for the Automatic Payload Deployment System*, Association for Unmanned Vehicle Systems International.
- Qin, Y., He, Z., Voigt, T., 2011. Towards Accurate and Agile Link Quality Estimation in Wireless Sensor Networks. *Ad Hoc Networking Workshop, 2011 the 10th IFIP Annual Mediterranean*, 179-185.
- Rondinone, M., Ansari, J., Riihijärvi, J., Mähönen, P., 2008. Designing a Reliable and Stable Link Quality Metric for Wireless Seonsor Networks. *REALWSN'08*.
- Srinivasan, K., Kazandjieva, M.A., Jain, M., Levis, P. 2008. PRR Is Not Enough.
- Srinivasan, K., Levis, P., 2006. RSSI Is Under Appreciated. *Third Workshop on Embedded Networked Sensors*.
- Vlavianos, A., Law, L.K., Broustis, I., Krishnamurthy, S.V., Faloutsos, M., 2008. Assessing Link Quality in IEEE 802.11 Wireless Networks: Which is the Right Metric? *IEEE 19th International Symposium on Personal, Indoor and Mobile Radio Communications*, 1-6.